
HIGH PERFORMANCE GPU CODE GENERATION FOR MATRIX-MATRIX MULTIPLICATION USING MLIR: SOME EARLY RESULTS

Navdeep Katel^{1,2}Vivek Khandelwal¹Uday Bondhugula^{1,2}

¹ Dept of CSA, Indian Institute of Science
Bengaluru, Karnataka 560012, India.

navdeepkatel@iisc.ac.in, vivekkhandel@iisc.ac.in, udayb@iisc.ac.in

² PolyMage Labs

Entrepreneurship Centre, Indian Institute of Science
Bengaluru, Karnataka 560012, India.

navdeep@polymagelabs.com, uday@polymagelabs.com

ABSTRACT

This report presents some early results on code generation targeting tensor cores on NVIDIA GPUs using the MLIR compiler infrastructure. The state-of-the-art in high-performance deep learning today is primarily driven by highly tuned libraries. These libraries are often hand-optimized and tuned by expert programmers using low-level abstractions with significant effort. A lot of this effort may have to be repeated for similar hardware and future ones. This process is thus not modular or reusable to the same extent that compiler infrastructure like LLVM are. Manual optimization typically does not use a standard intermediate representation (IR), although the optimizations performed can be encoded as a sequence of transformation steps and customized passes on an IR. Hand tuning may also miss exploration of design points only reachable easily by automatic code generation. We believe that until the recent introduction of MLIR [14, 16] (Multi-level intermediate representation), IR infrastructure was not geared to tackle the problem of automatic generation of domain-specific libraries in an effective manner. In particular, it was hard to represent and transform compute abstractions at high, middle, and low levels using a single IR.

With suitable abstractions in MLIR, we build an experimental lowering pipeline that is able to automatically generate code for matrix-matrix multiplication on NVIDIA GPUs targeting its tensor cores. On a set of problem sizes we evaluated, initial performance results show that we are able to attain performance that is 95-119% and 80-160% of CuBLAS [19] for F32 and F16 accumulate respectively on NVIDIA's Ampere [18] microarchitecture-based GeForce 3090 RTX. We believe that these results could be used as motivation for further research and development on automatic code and library generation using IR infrastructure for similar specialized accelerators.

1 Introduction

Deep learning and high-performance artificial intelligence in general heavily relies on high-performance computing today. The associated computing demands are being constantly met with innovations in computer hardware and micro-architecture, libraries, compilers, runtimes, and programming models. A significant amount of high-performance deep learning is currently powered by highly tuned libraries provided by hardware vendors: for example, CuDNN, CUBLAS, and MKL. Creating these libraries involves significant effort and expertise. This development process may have to be repeated with every major hardware or software release, and there are limits to what can be explored and optimized in an efficient manner.

Matrix-matrix multiplication (matmul) as a computational kernel is at the heart of many deep learning frameworks based on Transformers [26] like BERT [7], and high-performance computing in general. It also serves as an excellent test case to gauge what can be achieved. While automatic code generators’ strength is often optimizing a composition of kernels as opposed to a well-known individual one, not being able to generate good code for a well-studied kernel can be a showstopper for the overall story of having automatic code generators tackle all peak performance code generation. In this report, we specifically target tensor cores on NVIDIA GPUs, which are specialized units for Matrix-Multiply-Accumulate (MMA) operations with typically 3-4x more throughput than normal CUDA cores.

A few recent works have focused on GPU GEMM targeting tensor cores. Faingnaert et al. [8] try to address the *two language problem* by creating a three-level API in Julia [3] that enables the user to write efficient GEMM kernels. Their primary focus is on developing an API flexible enough to cater to a wide variety of applications, it is not with a unified IR infrastructure with multiple levels of abstraction. Bhaskaracharya et al. [4] take a polyhedral code generation approach to generate code for Volta tensor cores. They use schedule trees to represent the computation and use ISL [27] to generate CUDA code for it. They can generate code for matmul and fused operations like bias addition and ReLU, while achieving speed-ups up to $2.55\times$. The work is specific to Volta, and includes some device specific specializations to achieve competitive performance. Tillet et al. [25] recently presented Triton, an IR and an optimizing compiler for neural network computations. The framework is based on the concept of a *tile*, which is a statically shaped multi-dimensional array. The compiler is exposed by a python package which allows the user to write python code for which efficient machine code will be automatically generated. This work supports both CUDA and tensor core as well and achieves great performance.

Our approach here is to use compiler intermediate representation (IR) infrastructure to facilitate high-performance code and library generation. We use matrix-matrix multiplication as the kernel to experiment and NVIDIA’s tensor cores [2] as the target. MLIR [14, 15, 16] is the compiler infrastructure we use here with the goal to make the whole process more modular, systematic, and automatic to a large extent. We show that by lowering the IR in small steps and applying the right set of IR transformations and optimizations, we can achieve performance that is comparable to those of hand-written libraries without actually writing any code by hand. While a previous work [5] had done a similar study on a single core of a high-performance CPU, we are targeting a specialized accelerator here.

Some of the contributions of this work include:

- the introduction of Warp Matrix Multiply Accumulate (WMMA) [13] operations in an MLIR dialect along with their lowering to the LLVM/NVPTX backend,
- demonstrating how matmul on GPUs can be systematically and progressively code generated as a sequence of MLIR transformation and dialect lowering passes,
- building an end-to-end matmul code generation pipeline targeting tensor cores and preliminary results show that the performance obtained is on par with that of hand-tuned libraries and in some cases up to $1.60\times$ better.

Our approach being an IR-based one can work with different programming models and languages if a lowering path from such models to MLIR exists.

2 Background

2.1 MLIR

Multi-Level Intermediate Representation (MLIR) [14] aims to provide reusable, extensible compiler infrastructure and reduce the cost of building domain-specific compilers. MLIR can be used to serve multiple objectives such as:

- It can represent dataflow graphs (such as in TensorFlow), including dynamic shapes, variables, etc.
- It can be used to represent kernels for ML operations in a form suitable for optimization.
- It has the ability to host high-performance-computing-style loop optimizations across kernels (fusion, loop interchange, tiling, unroll-and-jam).
- It can represent target-specific operations, e.g., accelerator-specific high-level operations.
- It can represent kernels at different levels of abstractions (dialects in MLIR), which can help perform transformations and optimizations not possible at a single level.

Some recent works that employ MLIR include [9, 6].

The MLIR structure is made up of the following components:

- **Operations:** This is the basic unit of semantics in MLIR and is referred to as *Op*. Everything from instructions to functions to modules is modeled as Ops in MLIR. Ops take and produce zero or more *values*, called *operands* and *results* respectively.
- **Attributes:** It is structured compile-time static information e.g., integer constant values, string data, etc. Attributes are typed, and each Op instance has an open key-value dictionary from string names to attribute values.
- **Regions and blocks:** A *region* contains a list of blocks, and a block contains a list of operations that may further contain regions. The blocks inside the region make a Control Flow Graph(CFG). Each block ends with a *terminator* operation that may have *successor* blocks to which the control flow may be transferred.
- **Dialects:** It is a logical grouping of Ops, attributes, and types under a unique namespace. Ops from different dialects can coexist at any level of the IR at any time. Dialects allow extensibility and provide flexibility that helps in performing specific optimizations and transformations. Affine, GPU, LLVM, and Linalg are some of the important dialects.
- **Functions and modules:** A module is an Op with a single region containing a single block and terminated by a dummy Op that does not transfer the control flow. A function is an Op with a single region, with arguments corresponding to function arguments.

Some MLIR dialects that we have used in our work are explained below:

- **Affine Dialect:** This dialect uses techniques from polyhedral compilation to make dependence analysis and loop transformations efficient and reliable. We have performed most of the optimizations and transformations at the level of affine dialect.
- **GPU Dialect:** The GPU dialect models the general GPU programming paradigm similar to CUDA or OpenCL in MLIR. Its goal is to provide abstractions to model GPU-specific operations and properties. It is largely meant to be vendor-agnostic. Some additional information can be found in [11, 12] and the GPU dialect documentation [16].
- **NVVM Dialect:** Since we are focusing on tensor core code generation, we use and extend the NVVM dialect. This dialect provides operations that directly map to the NVPTX backend in LLVM.
- **LLVM Dialect:** The final stage of code generation involves lowering to LLVM IR, from where the LLVM backend takes control and generates the target code. To model LLVM IR, this dialect is used. This is the lowest level of abstraction present in MLIR.

2.2 GPU Background

GPUs are general-purpose massively parallel computing devices. The memory and the compute hierarchy play an essential role in optimizing any application and thus achieving high performance. We can abstract the GPU memories into a 4-level hierarchy, global memory, L2-cache, a configurable L1-cache/shared memory, and registers. The processors on GPUs can also be abstracted into a two-level hierarchy, the Streaming Multiprocessors (SMs) and computing cores inside the SM. The computing cores are often known as CUDA cores. Besides CUDA cores, special units like tensor cores are also present in newer GPUs at same level in the compute hierarchy. Each SM is further divided into processing blocks which have their respective warp schedulers. The programming model of GPUs is also structured to match the processor hierarchy present. A *thread* is a single execution entity on a GPU that can execute in parallel with other threads. These threads are clubbed into groups of 32, called a *warp*. Warps execute in a lockstep manner on the computing cores present on the SM. A *warp scheduler* chooses a warp that is ready to execute and dispatches it to the computing cores. When a warp encounters a data dependency, it stalls, and the warp scheduler chooses another warp that is ready for execution. Depending on the number of processing blocks present on the SM, multiple warps may be executing in parallel. So, in general, more warps help one achieve: (i) parallelism at the warp level,, (ii) better latency hiding,, and (iii) and better utilization of the underlying resources. Now, these warps are further grouped into a *thread block*. There can be several thread blocks that may be executing in parallel on the GPU. A thread block is bound to an SM. It cannot change the SM during its execution lifetime and must complete its execution on the same SM and release all the resources allocated to it upon completion. Threads in the same warp can exchange data using warp level shuffle instructions [28]. All the threads in a thread block can communicate using low latency *shared memory*, and threads in different thread blocks need to communicate using the high latency *global memory*. Synchronization primitives are present at thread block and warp level. A Synchronization will ensure that none of the threads in a thread block or a warp, depending on the type of synchronization used, proceed to the next instruction until all of them have arrived at the synchronization point. Synchronization becomes necessary while using shared memory in a case where data is first written to shared memory and then read by all threads. All threads must be synchronized before reading from and writing into the shared memory buffers to ensure correctness.

2.3 Tensor Cores

Tensor Cores [2] are programmable matrix-multiply-and-accumulate (MMA) units present on NVIDIA GPUs. First introduced in Volta architecture, they are present on Turing and Ampere as well. Significantly more throughput than CUDA cores makes them excellent for accelerating deep learning workloads. They perform MMA operation represented as, $\mathbf{D} = \mathbf{A} * \mathbf{B} + \mathbf{C}$, where the operations size is $4 \times 4 \times 4$ on Turing and Volta while it is $8 \times 4 \times 8$ on Ampere. Tensor cores execute *warp synchronous* instructions like HMMA to perform the MMA operation. Warp synchronous means that all the threads in a warp cooperatively execute these special instructions in order to produce the output of the MMA operation. Due to this warp synchronous nature of the tensor cores instructions, it becomes necessary to write or generate code at the warp level rather than the thread level when programming tensor cores. Tensor cores originally supported only FP16 for input and FP16 or FP32 for accumulation and output. But now they support a variety of formats for both input and output like TF32, BF16, INT8, and INT4. TF32 has the same range as FP32, and the same precision as FP16, but is represented in 19-bits. It can be used at places where precision is of less importance. To use this mode the inputs have to be in FP32, they will be internally converted to TF32, accumulation will done in FP32, and the output will be produced in FP32 as well. This provides a speed-up over normal FP32 mode on CUDA cores. BF16 offers same range as FP32, and precision which is less than FP16. Tensor cores offer the same speed in both BF16 and FP16 modes, while both are faster than TF32. The integer types are meant to be used in post training quantization [21].

When it comes to programmability, there are three ways to leverage tensor cores: (i) use a high level library like cuBLAS,, (ii) program using a high level C++ WMMA API [1] in CUDA, or, and (iii) program them explicitly using assembly level instructions.

	Performance	Shared memory bank conflicts	Ease of use	Support for operator fusion
High-level libraries	Best	minimal	Accessible via function calls	limited
WMMA API	Competitive in most cases	higher	Programming effort required	good
Assembly	May be competitive in all cases	minimal	Significant programmer effort required	good

Table 1: Comparison of different approaches to program tensor cores.

While cuBLAS can be used using mere functions calls, using the other two methods, require significant programming efforts. WMMA API provides larger matrix operations ($16 \times 16 \times 16$, $32 \times 8 \times 16$), and utility functions to load and store operand matrices. The task of converting these API functions to GPU microarchitecture specific assembly instructions is also offloaded to NVIDIA’s proprietary compile. The matrices loaded using WMMA API have and opaque layout once they are loaded into the registers, i.e., which thread is holding which element of the loaded matrix (thread-data mapping) is not known. Due to this opaque nature, some additional steps are needed when doing fusion with operations like bias-addition, which need to be aware of the thread-data mapping. Programming tensor cores explicitly using assembly instructions is even more challenging because the programmer has to take care of complexities such as thread data mappings in registers and data movement between shared memory and registers. The approaches are summarized in Table 1.

The NVPTX backend in LLVM exposes WMMA API functions as intrinsics. This makes it possible to program the tensor cores via MLIR. These intrinsics map one-to-one with the WMMA API functions, and exhibit identical behaviour in terms of programming and usage.

3 Design

In this section we present the design of our pipeline. Our pipeline is based on a series of optimizations and transformations which make up the recipe for fast matmul on GPUs. The recipe we use is close to what is highlighted in some prior works[[20],[8],[4]]. The part that is common in all of them is two-level blocking to maximize the reuse at different levels of the memory hierarchy. The general recipe is described in Algorithm 1.

Prior to our work, some support was present in MLIR, which we have reused in our pipeline, but some core components were missing. Mainly, the ops required to program the tensor cores using the WMMA API were absent from MLIR, and we were the ones to introduce these ops. We make changes and additions to the existing MLIR infrastructure wherever necessary.

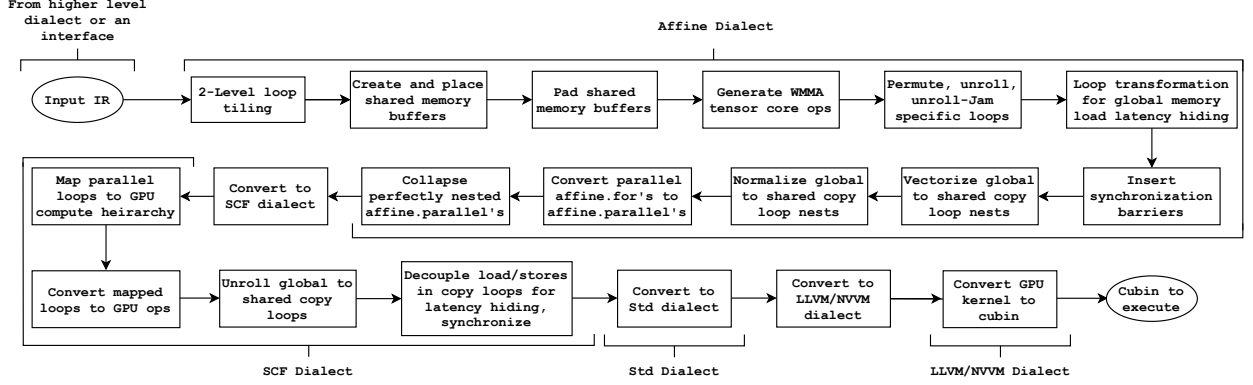


Figure 1: Lowering path for matmul to target GPU tensor cores.

Figure 1 shows the lowering path that we take, which is based on Algorithm 1. Although there can be different lowering paths to achieve the same goal, we argue that the one going through affine dialect should be chosen as long as the target kernel to be generated is affine. This could help in many things such as creation and placement of fast memory buffers, loop-tiling, unroll-jam, vectorization, detection of parallel loops, and placement of synchronization barriers too.

Algorithm 1: Two-level tiled tensor core Matmul.

```

Global memory: A[M][K] B[K][N], C[M][N];
Shared memory: a_smem[tbm][tbk], b_smem[tbk][tbn];
Registers: c_reg[wm][wn], a_reg[wm], b_reg[wn];
for threadBlockK  $\leftarrow$  0 to M step tbk do
  __syncthreads();
  All threads load  $tbm \times tbk$  block form A to a_smem;
  All threads load  $tbk \times tbn$  block form B to b_smem;
  Warp loads  $wm \times wn$  block form C to c_reg[wm][wn];
  __syncthreads();
  // wmmaM, wmmaN, and wmmaK represents the WMMA intrinsic size;
  for warpK  $\leftarrow$  0 to tbk step wmmaK do
    for warpM  $\leftarrow$  0 to wm step wmmaM do
      Warp loads a fragment of A form a_smem into a_reg[warpM];
      for warpN  $\leftarrow$  0 to wn step wmmaN do
        Warp loads a fragment of B form b_smem into b_reg[warpM];
        c_reg[warpM][warpN] += a_reg[warpM]  $\times$  b_reg[warpN];
      end
    end
  end
  Warp stores c_reg[wm][wn] to the respective block in C;
end

```

Although not highlighted in Algorithm 1 for brevity, It is worth noting that high performance can only be achieved using a set of some more optimizations, which include (i) padding in shared memory buffers to reduce bank conflicts, (ii) register tiling or warp tiling, (iii) load-store vectorization, and (iv) global memory load latency hiding. Now, we will describe our lowering pipeline in some detail, discussing how the major optimizations are enabled. We skip over the trivial parts.

3.1 A Starting Point

The starting point for our code generation approach is a high-level op like `lmhlo.dot` or `linalg.matmul` or simply an IR that is generated from a user-facing programming model targeting a linear algebra dialect. In the former case, we can lower the op to a three-loop affine matmul, while we can generate the three-loop affine matmul in the latter. Our starting point is shown in Listing 1.

```

affine.for %i = 0 to %M {
  affine.for %j = 0 to %N {
    affine.for %k = 0 to %K {
      %a = affine.load %A[%i, %k] : memref<8192x8192xf16>
      %b = affine.load %B[%k, %j] : memref<8192x8192xf16>
      %c = affine.load %C[%i, %j] : memref<8192x8192xf32>
      %aq = fpext %a : f16 to f32
      %bq = fpext %b : f16 to f32
      %q = mulf %aq, %bq : f32
      %co = addf %c, %q : f32
      affine.store %co, %C[%i, %j] : memref<8192x8192xf32>
    }
  }
}

```

Listing 1: Naive affine matmul.

3.2 Tiling for Locality and Parallelism

It is well known that tiling helps to exploit data reuse and significantly improves performance if the correct parameters are chosen for tiling. Two-level tiling is essential to achieve optimal performance in GPUs. The first level of tiles are mapped to different thread blocks. Each thread block copies its tile for **A** and **B** from global memory to shared memory, thus preventing multiple high latency trips to global memory. As the partitioned tiles are mapped to different thread blocks, they can be computed in parallel on different SMs. The second level of tiling is done to exploit reuse in the registers and promote warp-level parallelism. The thread block level tile is partitioned between the warps, and each warp only works on the part part of the tile mapped to it. This step gives us the two level tiled structure to proceed with.

3.3 Creating and Placing Shared Memory Buffers

After tiling is done, the next step is to create and place shared memory buffers at the right loop depth. We use `affineDataCopyGenerate` utility to generate the copies for matrices A and B. We take a slightly different approach here than some previous works. We create shared memory buffers only for matrices A and B. Since C is only loaded once per warp, we stream it directly into the registers from the global memory. Faingnaert et al [8] stream C from global memory to shared memory and then from shared memory to registers. The rationale for streaming C via shared memory is the prevention of random accesses to global memory and perhaps promoting coalesced accesses in global memory, which could be more efficient. We speculate that this may not always be the case, especially for large problem sizes, as C tile is only loaded once per warp.

Additionally, this approach would also require the use of dynamically allocated shared memory, as the optimal tile size for C could easily exhaust the 48 KB static shared memory limit on some devices. So the buffer to hold the tiles must be allocated dynamically and must be reused for storing tiles of all three operands. The support to allocate shared memory dynamically is currently absent from MLIR. So we restrict ourselves to statically allocated shared memory to avoid additional complexities in the code generator which may not be worth the effort. Even without doing this we already reach close to hand-tuned libraries in most cases.

Creating the shared memory buffers is one part of the story while ensuring shared memory accesses with minimal bank conflicts is another. Bank conflicts can reduce the throughput of shared memory by a significant amount. A general technique to avoid bank conflicts is to pad the shared memory buffers in the leading dimension. We achieve the same thing by changing the *leadingDimension* of the shared memory buffer that was generated by `affineDataCopyGenerate` to *leadingDimension + paddingFactor*. Doing this will change the underlying layout map of the shared memory buffer to account for the change in the leading dimension, and the rest of the IR need not be changed. It is also worth noting that we can try out different padding factors here and see what performs the best, but the padding factor must be a multiple of 8, i.e., 128-bits for F16 elements. This is because of an alignment requirement by the WMMA API.

3.4 Generating WMMA Ops and More

Now that we have all the basic things that we need, we can proceed towards the generation of `gpu.subgroup_mma` ops. The WMMA ops are of different sizes and we use the $16 \times 16 \times 16$ version of the ops in this work. The ops we

generate here should replace the scalar ops that are already present, and the loop steps of the corresponding loops must be adjusted accordingly.

```
#map0 = affine_map<(d0) -> (d0)>
#map1 = affine_map<(d0) -> (d0 + 64)>
#map2 = affine_map<(d0) -> (d0 + 128)>
module {
  // Shared memory buffers for A and B.
  memref.global "private" @b_smem_global : memref<64x136xf16, 3>
  memref.global "private" @a_smem_global : memref<128x72xf16, 3>
  func @main() {
    ...
    affine.for %i = 0 to 8192 step 128 {
      affine.for %j = 0 to 8192 step 128 {
        // References to shared memory buffers.
        %b_smem = memref.get_global @b_smem_global : memref<64x136xf16, 3>
        %a_smem = memref.get_global @a_smem_global : memref<128x72xf16, 3>
        // Main k-loop.
        affine.for %k = 0 to 8192 step 64 {
          // Copy loop for B.
          affine.for %copykk = #map0(%k) to #map1(%k) {
            affine.for %copyjj = #map0(%j) to #map2(%j) {
              %11 = affine.load %B[%copykk, %copyjj] : memref<8192x8192xf16>
              affine.store %11, %b_smem[%copykk - %k, %copyjj - %j] : memref<64x136xf16, 3>
            }
          }
          // Copy loop for A.
          affine.for %copyii = #map0(%i) to #map2(%i) {
            affine.for %copykk = #map0(%k) to #map1(%k) {
              %11 = affine.load %A[%copyii, %copykk] : memref<8192x8192xf16>
              affine.store %11, %a_smem[%copyii - %i, %copykk - %k] : memref<128x72xf16, 3>
            }
          }
          affine.for %ii = 0 to 128 step 64 {
            affine.for %jj = 0 to 128 step 32 {
              affine.for %kk = 0 to 64 step 32 {
                affine.for %kkk = 0 to 32 step 16 {
                  affine.for %iii = 0 to 64 step 16 {
                    affine.for %jjj = 0 to 32 step 16 {
                      ...
                      %a = gpu.subgroup_mma_load_matrix %a_smem[%11, %12] {leadDimension = 72 :
                        index} : memref<128x72xf16, 3> -> !gpu.mma_matrix<16x16xf16, "A0p">
                      %b = gpu.subgroup_mma_load_matrix %b_smem[%12, %14] {leadDimension = 136 :
                        index} : memref<64x136xf16, 3> -> !gpu.mma_matrix<16x16xf16, "B0p">
                      %c = gpu.subgroup_mma_load_matrix %C[%16, %17] {leadDimension = 8192 : index
                        } : memref<8192x8192xf32> -> !gpu.mma_matrix<16x16xf32, "C0p">
                      %res = gpu.subgroup_mma_compute %a, %b, %c : !gpu.mma_matrix<16x16xf16, "A0p
                        ">, !gpu.mma_matrix<16x16xf16, "B0p"> -> !gpu.mma_matrix<16x16xf32, "
                        C0p">
                      gpu.subgroup_mma_store_matrix %res, %C[%16, %17] {leadDimension = 8192 :
                        index} : !gpu.mma_matrix<16x16xf32, "C0p">, memref<8192x8192xf32>
                    }
                  }
                }
              }
            }
          }
        }
      }
    }
  }
}
```

Listing 2: Tiled and padded affine matmul with WMMA ops.

Now that we have generated the WMMA operations, we do the following IR transformations:

- permute the outermost six loops to go from (i, j, k, ii, jj, kk) order to (i, j, ii, jj, k, kk) order. This later helps in mapping compute loops to GPU compute hierarchy. Additionally, it also helps in moving invariant load-store operations on C to the most outermost position possible.
- permute the innermost three loops to go from (i, j, k) to (k, i, j). This represents warp level MMA operation as an outer product and enhances ILP, as pointed out by Bhaskaracharya et al. [4].
- fully unroll the innermost three loops.

Listing 2 shows the IR that we get after the creation of WMMA ops. We should note the step of the innermost loops being adjusted here. The listing further shows the loop nest in a permutation we desire. The two outermost loops will later be mapped to the thread blocks in a grid, and the following two loops will be mapped to the warps. The next two loops are the k-loops corresponding to the thread block and then the warp. After unrolling, we make two observations (i) the operations on C matrix now become independent of the two surrounding loops, and hence we hoist the operations on C now to the outermost k-loop. In this way, we prevent repetitive loads and stores to C in the global memory and only do them when the processing of a thread block tile begins and ends, and (ii) unrolling these loops completely reveal all loads on A and B. Some of these loads are the same along the k-dimension, and by applying CSE, we can completely remove the redundant loads and achieve unroll-jam kind of effect.

```

...
#map0 = affine_map<(d0, d1) -> (d0 + d1)>
...
// Thread block 'i' loop.
affine.for %i = 0 to 8192 step 128 {
  // Thread block 'j' loop.
  affine.for %j = 0 to 8192 step 128 {
    %b_smem = memref.get_global @b_smem_global : memref<64x136xf16, 3>
    %a_smem = memref.get_global @a_smem_global : memref<128x72xf16, 3>
    // Warp 'i' loop.
    affine.for %ii = 0 to 128 step 64 {
      // Warp 'j' loop.
      affine.for %jj = 0 to 128 step 32 {
        // Hoisted loads on C.
        %i1 = affine.apply #map0(%i, %ii)
        %i2 = affine.apply #map0(%j, %jj)
        %c_reg_0 = gpu.subgroup_mma_load_matrix %C[%i1, %i2] {leadDimension = 8192 : index} :
          memref<8192x8192xf32> -> !gpu.mma_matrix<16x16xf32, "C0p">
        ...
        // Main 'k'-loop with loaded C operand as iter_args.
        %res:8 = affine.for %k = 0 to 8192 step 64 iter_args(%c_in_0 = %c_reg_0, %c_in_1 = %
          c_reg_1 ...) -> (!gpu.mma_matrix<16x16xf32, "C0p">, !gpu.mma_matrix<16x16xf32, "C0p
          ">) {
          ...
          %a = gpu.subgroup_mma_load_matrix %a_smem[%ii, %c_in_0] {leadDimension = 72 : index} :
            memref<128x72xf16, 3> -> !gpu.mma_matrix<16x16xf16, "A0p">
          %b = gpu.subgroup_mma_load_matrix %b_smem[%c_in_0, %jj] {leadDimension = 136 : index} :
            memref<64x136xf16, 3> -> !gpu.mma_matrix<16x16xf16, "B0p">
          %c_res = gpu.subgroup_mma_compute %a, %b, %c_in_0 : !gpu.mma_matrix<16x16xf16, "A0p">,
            !gpu.mma_matrix<16x16xf16, "B0p"> -> !gpu.mma_matrix<16x16xf32, "C0p">
          ...
          // Main 'k'-loop yielding the results of the current iteration.
          affine.yield %i04, %i07 ... : !gpu.mma_matrix<16x16xf32, "C0p">, !gpu.mma_matrix<16
            x16xf32, "C0p">...
        }
        // Hoisted stores on C.
        gpu.subgroup_mma_store_matrix %res#0, %C[%i1, %i2] {leadDimension = 8192 : index} : !gpu.
          mma_matrix<16x16xf32, "C0p">, memref<8192x8192xf32>
        ...
      }
    }
  }
}

```

Listing 3: Affine matmul after loop unrolling and invariant load-store hoisting.

The loop structure just after the said optimizations is shown in Listing 3. It must be noted how the loop structure has changed after the movement of invariant load-store pairs of C matrix. The `affine.for` operation in line 20 represents the main-k loop, and is now modified to take the loaded C operands as loop `iter_args`. These will serve as accumulators for the multiplications happening in this loop. After every iteration, this loop yields accumulated products, which are passed as `iter_args` to the next iteration. These `iter_args` reside in the registers and are reused across different iterations of the main k-loop.

3.5 Global Memory Load Latency Hiding

With the introduction of `gpu.subgroup_mma` ops and some other optimizations in the previous section, we are moving towards the structure that we will have in the final IR. We focus on doing as many optimizations as possible without having any GPU-specific information in the affine dialect itself. In the IR we have currently, we cannot start computation until we have loaded the shared memory tiles of A and B. Global memory loads are one of the most expensive operations in terms of latency, so eliminating the long wait times on the operands is very important. We do so by splitting the main k-loop or the thread block k-loop by taking out the copy of A and B for iteration 0 and compute for iteration $n - 1$. The copy is placed just before the k-loop, and compute is placed just after it. The indices for the computation being performed in that loop also need to be shifted to move one iteration ahead. As a result, compute happens on data already available in shared memory with loads for the next iteration having been issued. We show the structure of the IR at this stage in Listing 4.

```
#map4 = affine_map<(d0) -> (d0)>
#map5 = affine_map<(d0) -> (d0 + 128)>
#map6 = affine_map<(d0) -> (d0 + 64)>
// Peeled copy loops for iteration 0 of k-loop.
affine.for %copyk = 0 to 64 {
  affine.for %copyj = #map4(%j) to #map5(%j) {
    %35 = affine.load %B[%copyk, %copyj] : memref<8192x8192xf16>
    affine.store %35, %b_smem[%copyk, %copyj - %j] : memref<64x136xf16, 3>
  }
}
affine.for %copyi = #map4(%i) to #map5(%i) {
  affine.for %copyk = 0 to 64 {
    %35 = affine.load %A[%copyi, %copyk] : memref<8192x8192xf16>
    affine.store %35, %a_smem[%copyi - %i, %copyk] : memref<128x72xf16, 3>
  }
}
// Main k-loop.
affine.for %k = 0 to 8128 step 64 {
  // Copy loops for iteration '%k + 1' of k-loop.
  affine.for %copyk = #map6(%k) to #map5(%k) {
    affine.for %copyj = #map4(%j) to #map5(%j) {
      %36 = affine.load %B[%copyk, %copyj] : memref<8192x8192xf16>
      affine.store %36, %b_smem[%copyk - %k - 64, %copyj - %j] : memref<64x136xf16, 3>
    }
  }
  affine.for %copyi = #map4(%i) to #map5(%i) {
    affine.for %copyk = #map6(%k) to #map5(%k) {
      %36 = affine.load %A[%copyi, %copyk] : memref<8192x8192xf16>
      affine.store %36, %a_smem[%copyi - %i, %copyk - %k - 64] : memref<128x72xf16, 3>
    }
  }
  affine.for %kk= 0 to 64 step 32 {
    ...
  }
}
// Peeled compute loop for the last iteration of k-loop.
affine.for %arg4 = 8128 to 8192 step 64 {
  ...
}
```

Listing 4: WMMA affine matmul with shifted k-loop.

While this lays the groundwork for latency hiding, to see this in action, we need to decouple the stores into shared memory from the loads on global memory for the copy loops inside the thread block k-loop. This is required both for

the correctness and functioning of the optimization. For this, we unroll the copy loops and delay stores to be the trailing operations in the outer k-loop. We delay this optimization to a further point in the pipeline as some GPU-specific information is required to enable it.

3.6 Inserting Synchronization Barriers

We have finished generating most of the parts in the IR, which may require synchronization barriers. The shared memory buffers will be read and written by all the threads in the thread block, so it is essential to synchronize before and after writing into these buffers. In general, this process can also be automated using memory-based dependence analysis. However, we place these synchronization barriers using the aforementioned static information about the copy loops for our current purposes.

3.7 Global to Shared Copy Vectorization

While latency hiding will do its bit, it cannot make the actual copy go any faster. It is well known that vector load store instructions [17] perform better than their scalar counterparts because it reduces the number of memory transactions and often results in better utilization of the available bandwidth.

We use the vectorization utilities already present in MLIRX [24]. We call this utility on the global to shared memory copies. We can try out different widths for the vector widths using this utility. We tried out 32, 64 and 128 bit wide vectors and found out the 128-bit wide vectors to work the best. Vectorized copy loops are shown in Listing 5.

```
...
#map4 = affine_map<(d0) -> (d0)>
#map5 = affine_map<(d0) -> (d0 + 128)>
#map6 = affine_map<(d0) -> (d0 + 64)>
...
// Cast operations for global memory memrefs.
%a_cast = memref.vector_cast %A : memref<8192x8192xf16> to memref<8192x1024xvector<8xf16>>
%b_cast = memref.vector_cast %B : memref<8192x8192xf16> to memref<8192x1024xvector<8xf16>>
// Cast operations for shared memory memrefs.
%b_smem_cast = memref.vector_cast %b_smem : memref<64x72xf16, 3> to memref<64x9xvector<8xf16>,
3>
%a_smem_cast = memref.vector_cast %a_smem : memref<128x72xf16, 3> to memref<128x9xvector<8xf16>,
3>
// Vectorized copy loops.
affine.for %copyk = #map6(%k) to #map5(%k) {
  affine.for %copyj = #map4(%j) to #map5(%j) step 8 {
    %i35 = affine.load %b_cast[%copyk, %copyj floordiv 8] : memref<8192x1024xvector<8xf16>>
    affine.store %i35, %b_smem_cast[%copyk - %k - 64, (%copyj - %j) floordiv 8] : memref<64
      x17xvector<8xf16>, 3>
  }
}
affine.for %copyi = #map4(%i) to #map5(%i) {
  affine.for %copyk = #map6(%k) to #map5(%k) step 8 {
    %i35 = affine.load %a_cast[%copyi, %copyk floordiv 8] : memref<8192x1024xvector<8xf16>>
    affine.store %i35, %a_smem_cast[%copyi - %i, (%copyk - %k) floordiv 8 - 8] : memref<128
      x9xvector<8xf16>, 3>
  }
}
```

Listing 5: Vectorized copy loops.

3.8 Extracting Parallel Loops.

This is the last step that we take while still in the affine dialect. We use `isLoopParallel` utility present in MLIR to find all the parallel loops and then convert them to parallel loops using `affineParallelize`. These parallel loops are later processed and mapped to the GPU processor hierarchy, while the sequential loops are the only ones which will remain in the kernel.

3.9 Mapping to GPU Compute Hierarchy

The previous step was the last one in the affine dialect, and we immediately convert to the SCF dialect after it. Starting in the SCF dialect, the first thing that we do is map the parallel loops to the GPU compute hierarchy. The existing utilities and passes in MLIR for mapping do not support mapping loops to individual warps, which is required in our case. We extend the utilities and passes to add this support for matmul. Ideally, we should generalize the passes and utilities used in this step to handle a wide variety of loop nests, we leave it as future work. It is worth pointing out at this point that we take all the measures necessary to ensure coalesced global memory accesses [10] which are crucial for effective bandwidth utilization and faster copies from global to shared memory. After mapping is completed, the two outermost loops will be converted to a `gpu.launch` op, the next two loops will be mapped to the warps and the remaining compute loops are actually sequential and remain as is.

3.10 Completing Latency Hiding

In Section 3.5, we described latency hiding and concluded that it is not complete until we decoupled the loads from stores. To achieve this without introducing any complexity in the code we first completely unroll the copy loops inside the thread block k-loop and then delay the stores such that they happen after compute is completed. We show the general structure of the IR that we have in Listing 6. The approach we follow is quite similar to the one pointed out in [4].

```

gpu.launch blocks(%blockIdX, %blockIdY, %blockIdZ) in (%arg6 = %c64, %arg7 = %c64, %arg8 = %c1)
  threads(%threadIdX, %threadIdY, %threadIdZ) in (%arg9 = %c256, %arg10 = %c1, %arg11 = %c1)
  {
    ...
    %c_reg_0 = gpu.subgroup_mma_load_matrix %C[%26, %27] {leadDimension = 8192 : index} : memref
      <8192x8192xf32> -> !gpu.mma_matrix<16x16xf32, "C0p">
    ...
    // Peeled copy loops for iteration 0 of k-loop.
    scf.for %copy = %c0 to %c4 step %c1 {
      ...
    }
    scf.for %copy = %c0 to %c4 step %c1 {
      ...
    }
    gpu.barrier
    // Main k-loop
    %res:8 = scf.for %k = %c0 to %c8128 step %c64 iter_args(%c_in_0 = %c_reg_0, %c_in_1 = %c_reg_1
      ...) -> (!gpu.mma_matrix<16x16xf32, "C0p">, !gpu.mma_matrix<16x16xf32, "C0p"> ...) {
      gpu.barrier
      // Global memory loads for iteration i + 1 of k-loop
      %a_next_iter_0 = memref.load %a_cast[%74, %81] : memref<8192x1024xvector<8xf16>>
      %b_next_iter_0 = memref.load %b_cast[%94, %101] : memref<8192x1024xvector<8xf16>>
      ...
      scf.for %kk = %c0 to %c64 step %c32 iter_args(%arg16 = %c_in_0, %arg17 = %c_in_1) -> (!gpu.
        mma_matrix<16x16xf32, "C0p">, !gpu.mma_matrix<16x16xf32, "C0p"> {
        ...
      }
      gpu.barrier
      // Shared memory stores for iteration i + 1 of k-loop
      memref.store %b_next_iter_0, %b_smem_cast[%51, %68] : memref<64x17xvector<8xf16>, 3>
      memref.store %a_next_iter_0, %a_smem_cast[%150, %167] : memref<128x9xvector<8xf16>, 3>
      ...
    }
    gpu.barrier
    // Peeled compute loop for iteration n-1 of k-loop.
    scf.for %arg14 = %c0 to %c64 step %c32 {
      ...
    }
    gpu.subgroup_mma_store_matrix %res#0, %C[%26, %27] {leadDimension = 8192 : index} : !gpu.
      mma_matrix<16x16xf32, "C0p">, memref<8192x8192xf32>
    ...
  }

```

Listing 6: GobaI memory load latency hiding.

This is the terminal point for us in terms of optimizations, and this was our last step while in the SCF dialect.

3.11 Putting It All Together

As the previous step was the last one in terms of optimization, the generated IR is now to be set up for execution. There already exists well-designed infrastructure in MLIR to achieve this [12, 11] and we provide a summary outline of it for completeness. The existing design in MLIR allows one to represent IR to be executed on an accelerator like the GPU within a single MLIR file. The IR would have two parts: a host-side component, which runs on the CPU, and a device side component, or the kernel, which runs on the GPU. The host side component invokes the device side component, and can wait for its execution to complete or can proceed to do other tasks. The lowering paths are slightly different for the host and device side components:

- **Host side compilation:** The host side code is converted to the `std` dialect and then subsequently to the `llvm` dialect. During conversion to the `llvm` dialect, operations from the GPU dialect like `gpu.launch` are lowered to function calls to the CUDA driver and CUDA runtime API, through a thin wrapper interface provided by MLIR’s CUDA runtime library. MLIR is then translated to LLVM IR, and target code is generated. The execution of IR is then made possible through `mlir-cpu-runner` (using the MLIR’s LLVM Orc-based JIT). It takes as arguments the shared libraries to link against, where we can supply the library corresponding to the CUDA driver API.
- **Device side compilation:** The device side code is also converted to the `std` dialect and then to a mix of `llvm` and `nvvm` dialect. This is in turn translated to LLVM IR and then to PTX by the NVPTX backend in LLVM. PTX is then converted to *cubin* (CUDA binary format) using NVIDIA’s compiler. NVIDIA’s compiler is called via the CUDA driver API from MLIR. The `gpu-to-cubin` pass in MLIR has access to the driver API and performs the PTX to cubin compilation and assembly for us. We extended the pass to take this pass to take additional options like optimization level and the maximum number of registers per thread, which are required while compiling PTX to cubin.

The infrastructure to perform these final steps already existed in MLIR. While our evaluation used the MLIR JIT, ahead-of-time compilation could also be performed with a similar setup. The GPU dialect operations that we introduced, `gpu.subgroup_mma_load_matrix`, `gpu.subgroup_mma_store_matrix` and `gpu.subgroup_mma_compute`, were open-sourced and upstreamed to the official LLVM/MLIR repository.

4 Evaluation

In this section, we present the performance of our kernels and compare them with CuBLAS 11.2. The evaluation is performed on an NVIDIA Ampere-based Geforce RTX 3090 installed on a x86-64 system with an AMD Ryzen Threadripper 3970X CPU, running Ubuntu 20.04 LTS. The following parameters were set for all the experiments:

- The SM clocks were set to the boost frequency mentioned in the whitepaper for all the experiments, which is 1695 MHz.
- We limit ourselves to statically allocated shared memory, which is equal to 48 KB.
- The maximum number of registers per thread is set to 255.

We use NVIDIA Nsight Systems [23] for timing and consider only the kernel runtime for calculating the attained TFLOPs. This applies to our kernels as well as CuBLAS. We consider different combinations of thread block level tiles and warp level tiles and report the best performing version. The reported performance has been averaged over ten runs.

Finally, we consider a matmul of the form $C = AB + C$ (all three matrices are stored in a row-major layout). We use the `m16n16k16` version of the WMMA intrinsic and limit ourselves to square problem sizes ranging from 1024 to 16384 with a step of 256. We assume that the problem sizes are multiples of thread block tiles, which are again multiples of warp tiles, which are in turn multiples of the WMMA intrinsic.

4.1 Mixed Precision Performance

In this section, we present the performance of our automatically generated mixed precision kernels. Matrix-matrix multiplication having A, B in F16 and the accumulation of the products being done in F32 is termed as mixed precision matmul. The output matrix C, is also in F32. This version is of particular interest to us because of its importance in training deep learning models [22].




Figure 2: Mixed precision (fp16 input, fp32 accumulation and output) performance on square sized matrices.

We achieve performance that is consistently within 95-119% of CuBLAS 11.2. Comparing our performance with the absolute peak of the device, we sustain 95.4% of the device peak. Figure 2 shows the performance of our automatically generated kernels on Ampere RTX 3090. The figure shows we are very close to CuBLAS on the evaluated sizes. On a few smaller sizes, we outperform CuBLAS. In general, CuBLAS kernels may not be as well-tuned for smaller sizes as they are for the larger ones. On larger sizes MLIR-generated code is within 2-8% of cuBLAS performance. It was further observed that smaller thread block tile sizes like $64 \times 64 \times 64$ performed better on smaller problem sizes. This indicates that smaller problem sizes are benefited by increased occupancy. While smaller tile sizes reduce reuse for A and B in shared memory, they increase occupancy and this benefits small problem sizes that may launch relatively fewer thread blocks. On large problem sizes, small tile sizes result in a relatively large number of small threads blocks, and the effect of reduced data reuse becomes more prominent, and increased occupancy does not benefit any more.

An automatic code generation approach also allows us to study the impact of individual optimizations by enabling or disabling optimizations selectively. We show the impact of each of the optimizations discussed earlier in Figure 3 in an incremental manner starting from a naive version to the fully optimized one.

4.2 Half Precision Performance

In this section, we present the performance of our automatically generated half precision kernels. In this version of matmul, all three matrices A, B and C are in FP16. The accumulation of products is also done here in FP16. This version is in general faster than F32 version, but can be prone to imprecision due to a narrower representation for the mantissa and exponent.

We achieve performance that is consistently 80-160% of cuBLAS 11.2. Figure 4 shows the performance of our automatically generated kernels on Ampere RTX 3090. We observe that cuBLAS has inconsistent performance throughout the range, particularly on problem sizes larger than $W = 8848$. This suggests that cuBLAS is not well-tuned for all problem sizes. In particular, on profiling cuBLAS kernels we observe that thread block tile sizes chosen by cuBLAS were actually smaller than the ones for which we had optimal performance, for example, for $W=11264$, cuBLAS chooses $128 \times 128 \times 32$, while we choose $128 \times 256 \times 32$. We have a single stage of pipelining for hiding the latency of global memory loads while cuBLAS is using five stages. However, stalls on global memory loads were much more for cuBLAS. This could be as a result of sub-optimal latency hiding.




Figure 3: Mixed precision performance (fp16 input, fp32 accumulate and output) for $M = N = K = 8192$ with various optimizations enabled and disabled.

5 Conclusion And Future Work

We presented early results on automatic code generation targeting specialized matmul instructions supported by NVIDIA’s tensor cores. These preliminary results demonstrate that an automatic code-generator can achieve performance competitive to that of hand-tuned libraries in many cases. Experimental results on an NVIDIA GeForce 3090 RTX (based on the NVIDIA Ampere architecture) demonstrated the effectiveness of the presented approach. These results are only meant to serve as a stepping stone for the design of robust code and library generators that are able to not only optimize a single kernel but enable composition and fusion of kernels. This is an area where it is well-known that optimized libraries have limitations. While there have been many efforts to allow fusion and code generation via DSL compilers or graph rewriters, a robust approach based on a unified IR infrastructure has still been missing.

6 Acknowledgements

We thank the MLIR community for the excellent infrastructure that this work rests on. We especially thank the developers of Affine, GPU, and LLVM dialects for providing useful infrastructure that was used extensively in this work. We also thank other members of the community who reviewed artifacts of this work that were already upstreamed to the official LLVM/MLIR repository.

References

- [1] WMMA API for CUDA. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#wmma>, 2021.
- [2] J. Appleyard and S. Yokim. Programming tensor cores in CUDA 9. <https://devblogs.nvidia.com/programming-tensor-cores-cuda-9/>, 2017.
- [3] J. Bezanson, A. Edelman, S. Karpinski, and V. B. Shah. Julia: A fresh approach to numerical computing. *SIAM Review*, 59(1):65–98, 2017.




Figure 4: Performance with fp16 (input, accumulation and output in fp16) on square sized matrices.

- [4] S. G. Bhaskaracharya, J. Demouth, and V. Grover. Automatic kernel generation for volta tensor cores. *CoRR*, abs/2006.12645, 2020.
- [5] U. Bondhugula. High performance code generation in MLIR: An early case study with GEMM. *CoRR*, 2003.00532, 2020. <https://arxiv.org/abs/2003.00532>.
- [6] L. Chelini, A. Drebes, A. Zinenko, A. Cohen, N. Vasilache, T. Grosser, and H. Corporaal. Progressive raising in multi-level ir. In *International Conference on Code Generation and Optimization (CGO)*, February 27th - March 3rd, 2021, Virtual Conference, 2021.
- [7] J. Devlin, M. Chang, K. Lee, and K. Toutanova. BERT: pre-training of deep bidirectional transformers for language understanding. *CoRR*, abs/1810.04805, 2018.
- [8] T. Faingnaert, T. Besard, and B. D. Sutter. Flexible performant GEMM kernels on gpus. *CoRR*, abs/2009.12263, 2020.
- [9] T. Gysi, C. Müller, O. Zinenko, S. Herhut, E. Davis, T. Wicky, O. Fuhrer, T. Hoeftler, and T. Grosser. Domain-specific multi-level IR rewriting for GPU. *CoRR*, abs/2005.13014, 2020.
- [10] M. Harris. How to Access Global Memory Efficiently in CUDA C/C++ Kernels. <https://developer.nvidia.com/blog/how-access-global-memory-efficiently-cuda-c-kernels/>, 2013.
- [11] S. Herhut. MLIR on GPUs, Apr. 2020. MLIR Open Design Meeting, Apr 16, 2020.
- [12] S. Herhut and O. Zinenko. GPUs in MLIR, Dec. 2019. MLIR Open Design Meeting, Dec 12, 2019.
- [13] N. Katel. Introduction of WMMA ops for tensor cores in MLIR. <https://reviews.llvm.org/rG875eb523c13249114507cb8facd797773e278d9e>, 2021.
- [14] C. Lattner, M. Amini, U. Bondhugula, A. Cohen, A. Davis, J. Pienaar, R. Riddle, T. Shpeisman, N. Vasilache, and O. Zinenko. MLIR: A compiler infrastructure for the end of Moore’s law. *CoRR*, abs/2002.11054, 2020. <http://arxiv.org/abs/2002.11054>.
- [15] C. Lattner, M. Amini, U. Bondhugula, A. Cohen, A. Davis, J. Pienaar, R. Riddle, T. Shpeisman, N. Vasilache, and O. Zinenko. Mlir: Scaling compiler infrastructure for domain specific computation. In *2021 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*, pages 2–14, mar 2021.

- [16] LLVM/MLIR. MLIR: A multi-level intermediate representation, 2020. <https://mlir.llvm.org>.
- [17] J. Luitjens. Increase Performance with Vectorized Memory Access. <https://developer.nvidia.com/blog/using-shared-memory-cuda-cc/>, 2013.
- [18] NVIDIA. Nvidia Ampere GA102 GPU architecture. <https://images.nvidia.com/aem-dam/en-zz/Solutions/geforce/ampere/pdf/NVIDIA-ampere-GA102-GPU-Architecture-Whitepaper-V1.pdf>, 2020.
- [19] NVIDIA. CuBLAS. <https://docs.nvidia.com/cuda/cublas>, 2021.
- [20] NVIDIA. CUTLASS. <https://github.com/NVIDIA/cutlass>, 2021.
- [21] NVIDIA. Improving INT8 Accuracy Using Quantization Aware Training and the NVIDIA Transfer Learning Toolkit. <https://developer.nvidia.com/blog/improving-int8-accuracy-using-quantization-aware-training-and-the-transfer-learning-toolkit/>, 2021.
- [22] NVIDIA. Mixed precision training in deep learning. <https://docs.nvidia.com/deeplearning/performance/mixed-precision-training/index.html>, 2021.
- [23] NVIDIA. Nsight Systems. <https://docs.nvidia.com/nsight-systems/UserGuide/index.html>, 2021.
- [24] PolyMage Labs. MLIRX. <https://github.com/polymage-labs/mlirx>, 2020.
- [25] P. Tillet, H. T. Kung, and D. Cox. Triton: An Intermediate Language and Compiler for Tiled Neural Network Computations. In *Proceedings of the 3rd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages*. Association for Computing Machinery, 2019.
- [26] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. *CoRR*, abs/1706.03762, 2017.
- [27] S. Verdoolaege. isl: An integer set library for the polyhedral model. In K. Fukuda, J. v. d. Hoeven, M. Joswig, and N. Takayama, editors, *Mathematical Software – ICMS 2010*, pages 299–302, Berlin, Heidelberg, 2010. Springer Berlin Heidelberg.
- [28] V. G. Yuan Lin. Using CUDA Warp-Level Primitives. <https://developer.nvidia.com/blog/using-cuda-warp-level-primitives/>, 2018.